

Cursor for Everyone: Follow Along

Cursor for Everyone: Follow Along

Turn a messy leadership call into a working dashboard **and** a board-ready slide deck. No coding background needed.

```
flowchart LR
  Listen["Listen to the call"] --> Structure["Structure the work"]
  Structure --> Build["Build the dashboard"]
  Build --> Verify["Verify it"]
  Verify --> Present["Present with slides"]
  Present --> Automate["Automate the loop"]
  Automate --> Listen
```

The loop: listen → structure → build → verify → present → automate. Today you do it once by hand. By the end, you’ll see how to make it repeat on its own.

Your role today

For the next hour, you’re on the team that owns a company’s internal **revenue dashboard**. A leadership call just landed a familiar ask in your lap.

Your goal	Ship a dashboard feature and a board-ready slide deck
Why it matters	This is everyday work: turning vague requests into real results

Get set up (about 5 minutes)

1. Install Cursor and sign in

Download Cursor from cursor.com if you haven’t already. Open it and sign in.

2. Get the workshop folder

No GitHub account, no Git, no command line. Pick whichever feels easier.

Option A (recommended): let Cursor's Agent grab it for you

1. In Cursor, open a new chat (the Agent panel, usually on the right).
2. Paste this and hit send:


```
Download this workshop and open it for me, no git needed. Use curl  
to fetch the ZIP,  
unzip it into my home folder, then open the unzipped folder:  
https://github.com/hsaab/cursor-for-everyone-  
workshop/archive/refs/heads/main.zip
```
3. When Cursor asks to run a command, click **Run**. It downloads the files, unzips them, and opens the folder.

That is the whole point of Cursor: you ask in plain English and the agent does the work. You just did your first one.

Option B: download it yourself (3 clicks)

1. Open [the workshop on GitHub](#).
2. Click the green **Code** button, then **Download ZIP**.
3. Unzip it. In Cursor, choose **File** → **Open Folder** and pick the unzipped folder.

What is GitHub? It's where this project lives online, like a shared drive for a project's files. This repo is public, so anyone can grab it.

Stuck on any step? Paste what you see (or any error) into the Cursor chat and ask it to help. That trick works for the rest of the workshop too.

3. Open the dashboard

Double-click **dashboard.html** in the project folder. It opens in your browser.

Notice the **“Revenue by GPU type”** card is empty. That's what we're going to fix.

Tip: In chat, type @ to point Cursor at a specific file (for example @transcript.md).

Connected tools are optional

Cursor can connect to tools you already use (Figma, Slack, Linear, and more) through **MCP** servers and **Plugins**. For today, **everything you need is already in this folder**. No extra accounts or setup.

What you need	In this repo	Optional live version
Call notes	transcript.md	Granola MCP

What you need	In this repo	Optional live version
Design	design/ folder	Figma MCP
Data	data/revenue_by_gpu_type.csv	Database MCP

Your host may demo the live versions on screen. You can try them later if you connect those tools. The repo always works on its own.

Step 1 · Understand the story

Before touching anything, read the call notes.

Open **transcript.md** in Cursor (click it in the left sidebar).

What's happening?

- Leadership loves the topline revenue number (~\$415M in June).
- But every review, someone asks: **“Which GPU types drive the revenue?”**
- The dashboard only shows totals and regional splits. Casey is manually building charts in spreadsheets every month.
- Riley wants to see the **mix shift over 6 months** (A100 tapering off, GB200 / Blackwell ramping up).
- Jordan needs **branded slides for the board** by Thursday.

The data already exists in data/revenue_by_gpu_type.csv. It just isn't in the dashboard yet.

That's your job today.

Step 2 · Understand what we're working with

The project files

File / folder	What it is
dashboard.html	The dashboard page (opens in your browser)
dashboard.js	Logic that draws the charts
data.js	The numbers baked into the page
styles.css	Colors and layout
data/	Raw data files (CSV spreadsheets)
design/	Design specs, colors, and mockups
.cursor/rules/	Standing instructions Cursor always follows

Take a minute to click through these in the sidebar. You don't need to understand the code. Just know where things live.

See the dashboard now

Open `dashboard.html` in your browser if it's not already open. Scroll through it. Notice what's there (totals, regional split, revenue trend) and what's missing (GPU breakdown).

Pick the right model

Cursor lets you choose which AI model powers your chat. Different models are good at different things.

Task	Model to use	Why
Planning, reasoning, reading messy docs	Opus (or similar reasoning model)	Better at understanding ambiguity and tradeoffs
Writing and editing code	Composer 2.5	Fast, cost-efficient, built for coding tasks
Quick questions about files	Any model works	Ask mode is read-only anyway

Cost tip: Opus is great for planning but can cost roughly **10x more** than Composer 2.5 for similar implementation work. A good pattern: **plan with Opus, build with Composer 2.5.**

Learn more: [Models and pricing](#)

Know your modes

Cursor has four chat modes. Press **Shift + Tab** to switch between them, or use the dropdown in the chat panel.

Mode	Best for	Can edit files?
Agent	Building features, making changes	Yes
Plan	Complex work where you want to review the approach first	Yes (after you approve)
Ask	Understanding code, exploring without changes	No (read-only)
Debug	Tricky bugs that need investigation	Yes

Learn more: [Agent overview](#)

Step 3 · Turn the messy call into a spec (Plan mode)

Now let's structure the work before building anything.

1. Switch to **Plan** mode (Shift + Tab until you see "Plan").

2. Select a **reasoning model** like Opus for this step.
3. Paste this prompt:

```
@transcript.md Read this internal revenue review call and list exactly what they're asking for as a short, concrete spec. Group it into "dashboard changes" and "slides". Keep it to bullet points.
```

Cursor will read the transcript, ask clarifying questions if needed, and produce a structured plan you can review and edit before any code gets written.

That's the power of Plan mode: think first, build second.

Step 4 · Iterate on the plan

Your plan is a draft. Now make it sharper by adding the design and teaching Cursor your preferences.

Add the design

Tell Cursor about the mockup so the build matches what the designer intended. In the same Plan chat (or a new one), paste:

```
@design/gpu-section-spec.md @design/design-tokens.json
@design/revenue-by-gpu-type.png
Review the plan and update it to include building this GPU section to match this design spec and these color tokens.
```

Set house rules with Rules

Rules are standing instructions Cursor follows automatically. You write them once, and every future chat benefits.

This repo already has two in `.cursor/rules/`:

Rule	What it does
<code>cursor-brand.mdc</code>	Cursor colors and clean visual style (why the dashboard and slides stay on-brand)
<code>dashboard-html.mdc</code>	How to format the dashboard's HTML and CSS

Open one in the sidebar. They're short.

Now create your own. Type `/create-rule` in chat and describe what Cursor should always do:

/create-rule Every chart on the dashboard needs labeled axes and a one-line

"so what?" caption underneath. Apply it to *.html.

Cursor saves it to .cursor/rules/ and follows it from then on.

Step 5 · Build it (Agent mode)

Time to ship. Switch to **Agent** mode and select **Composer 2.5**.

Paste this prompt:

```
@data/revenue_by_gpu_type.csv @design/gpu-section-spec.md
@design/design-tokens.json @dashboard.js @data.js @styles.css
Add a "Revenue by GPU type" section to the dashboard using this CSV,
matching the
design spec and tokens. Read the CSV and inline the data into the app
(do NOT fetch
it at runtime, since the page is opened directly from a file). Show
revenue by GPU
type for the latest month, plus a 6-month trend that makes the mix
shift obvious.
Use the GPU-type colors from the tokens. Call out the fastest-growing
GPU type and
flag where GB200 overtakes A100.
```

What happens when you hit send

The Agent will:

1. **Read** the files you tagged with @
2. **Search** the codebase for patterns to follow
3. **Edit** multiple files (data, HTML, CSS, JavaScript)
4. **Show you** what changed so you can review

You don't write code. You describe what you want, point at the right files, and the agent figures out the rest.

Step 6 · Verify the dashboard (your first Skill)

Two ways to check your work: let a skill do it, then trust your own eyes.

Run a ready-made Skill

Skills are reusable workflows you run by name, like saved recipes. Describe the steps once, and anyone can trigger them with a slash command.

This repo ships one for exactly this moment: **/verify-dashboard**. It reads your files and checks the new section against the workshop's acceptance criteria. Type it in chat:

```
/verify-dashboard
```

It reports a short pass/fail list: is the card there? does June total \$415M? is the GB200-over-A100 crossover shown? is the data inlined instead of fetched? If anything failed, paste the suggested fix and run it again.

Look with your own eyes

Go back to your browser and **refresh** `dashboard.html`. Check the “**Revenue by GPU type**” card:

- Latest month breakdown by GPU type
- 6-month trend showing the mix shift
- GB200 highlighted as the fastest grower
- The moment GB200 overtakes A100 called out

You just shipped a feature, and verified it two ways.

If you're ahead, try one of these follow-ups in Agent mode:

Change the GPU-type chart to a stacked view so I can see the mix each month.

Recolor the chart so GB200 stands out and the legacy A100 is muted.

Play around. This is your dashboard now.



Step 7 · Make your own Skill (and your slides)

You just *ran* a skill. Now *build* one. First, knock out the deck you owe leadership with a second ready-made skill.

Run the slides Skill

This repo also ships **/branded-deck**. It reads your dashboard, pulls the real numbers, applies the brand rule, and generates a polished `.pptx`. No Slack, Linear, Notion, or internal tools needed, just the files in this repo.

Type **/branded-deck** in chat and paste:

```
/branded-deck Summarize this for leadership: headline revenue, revenue by GPU type, the 6-month mix shift toward GB200, and 3 takeaways. Read the numbers from @dashboard.html and @data/revenue_by_gpu_type.csv.
```

Behind the scenes it:

1. Follows `.cursor/rules/cursor-brand.mdc` for on-brand styling
2. Reads the built dashboard and CSV for real numbers
3. Uses the built-in slide generator to produce the `.pptx`

Now create your own

You've run two skills (`/verify-dashboard` and `/branded-deck`). Making your own is just as easy. Type `/create-skill` and describe a workflow you repeat, for example:

```
/create-skill A "monthly-refresh" skill that re-reads
data/revenue_by_gpu_type.csv,
updates the dashboard numbers, runs /verify-dashboard, then runs
/branded-deck.
```

Cursor saves it under `.cursor/skills/` so you and your teammates can run it by name from then on.

⚡ Step 8 · Automate the whole loop (high level)

Everything you did today was manual: you pasted prompts, reviewed plans, and refreshed the browser. What if it happened automatically every month?

Cursor Automations are always-on cloud agents. You set a trigger (like a schedule), write instructions, and Cursor runs the job in the background without you in the IDE.

Here's what a monthly version of today's work could look like:

flowchart TD

```
Trigger["Monthly schedule fires"] --> FetchData["Agent grabs fresh
GPU revenue data"]
FetchData --> UpdateDash["Agent updates data.js and dashboard.html"]
UpdateDash --> Verify["Agent runs /verify-dashboard"]
Verify --> GenSlides["Agent runs /branded-deck"]
GenSlides --> Notify["Agent posts the deck for review"]
```

Step	What the automation does
Trigger	Runs on the 1st of every month (cron schedule)
Fetch data	Pulls the latest revenue-by-GPU export (from a CSV, database, or API)
Update dashboard	Inline the new numbers into <code>data.js</code> and refreshes the charts
Verify	Runs <code>/verify-dashboard</code> to catch data or rendering issues before slides
Generate slides	Runs the <code>/branded-deck</code> skill with the updated dashboard
Notify	Posts the new deck somewhere your team can review (email, Slack, or a shared

Step	What the automation does
	folder)

You wouldn't build this today. But now you know the shape: **the same loop you just did by hand, running on autopilot.**

Create one at cursor.com/automations or in the Agents window when you're ready.

Step 9 • Review what we did

Take a breath. Look at what you built.

Step	What you learned
1. Story	Read messy input and understand the real ask
2. Codebase	Know where files live and which model/mode to pick
3. Plan	Turn a call transcript into a structured spec (Plan mode)
4. Iterate	Add designs with @, set standing preferences with Rules
5. Build	Ship a feature with Agent mode + Composer 2.5
6. Verify	Run a ready-made Skill (/verify-dashboard), then eyeball it
7. Present	Run the /branded-deck skill, then build your own with /create-skill
8. Automate	See how the whole loop could run on a schedule (Automations)

What Cursor learned from you:

- A **rule** (your chart caption preference) that applies to every future chat
- A **skill** of your own, on top of the ready-made /verify-dashboard and /branded-deck, runnable by name anytime

What you can do next:

Swap in *your* call notes, *your* spreadsheet, *your* design files. The loop stays the same: listen → structure → build → verify → present → automate.

That's Cursor for everyone.